

Yulun Zeng

When I first started, I was planning to do multiple iterations of dead code elimination, constant propagation, and constant folding. When I got that working, I wanted to try local value numbering.

When I do local value numbering, I spent a lot of time deciding which data structure to use. I did not hash the instruction at first, rather I focused on copy propagation and constant folding. But at some point, I was losing high level instructions even when I comment out the existing optimization. The code also became so buggy that I could proceed no further.

Then I decided to get rid of local value numbering, rather try register allocation only. I used a map to keep track of which register was used by which operand, but it didn't work well with memory references. When I encounter memory references, I started getting errors like

unsupported instruction `mov`

and

incorrect register `%r13` used with `l` suffix

I believe this problem as a result of poor communication between lowlevel codegen and register allocation, but at this time I really don't have any late hours left to fix anything.

I decided to go with my one and only working constant propagation + constant folding.

I use a map that keeps track of the constant value of registers and their constant values. But before that, there are several conditions to check:

There are 2 types of instructions, the ones with a destination and the ones without.

If we further divide these two groups, there are 2 operands instructions like add, sub... and 1 operand instructions like move, localaddr...

And when we go further, the destination operand can be either memory reference, a plain virtual register with base value.

We should definitely not propagate constants to memory references and registers ranging from vr0 to vr9, but we can replace other virtual registers with constants if such mapping exists.

Now we know which types of virtual registers can be replaced, we need to know when can we replace.

We must check the liveness of registers, if they are alive after a basic block, we don't want to replace them, because they might have been assigned a value that will be used later, such as a loop variable. If we replace those, the loop value will be stuck.

Next is updating or deleting the mapping. After each assignment, for example `mov vr10, $3`, the mapping of vr10 to constant will be updated. But when it's a `mov vr10 vr11`, we will delete the mapping for the sake of constant propagation (In my failed local value numbering code, I tried to keep a `map<int, Operand>`, but it didn't work well for memory references).

If we arrive at an add/sub/div/mul operation with 2 constant operands, we can perform a calculation, update the destination mapping, and ignore the instruction.

The next thing is do is repeat the constant propagation and constant folding for 10 times, this method effectively reduced roughly 12% lowlevel code for example 29 and 31, and the execution time on average shortened from ~1.8s to ~1.7s. It could be helpful for programs with a lot of constant calculations, but not so much for dynamic values.

Some snippets of before and after for the high level codes:

Before:

```
mov_l   vr16, $500
```

```
mov_l   vr10, vr16
```

```
localaddr vr16, $0
```

```

mov_l   vr17, $0
sconv_lq vr18, vr17
mul_q   vr18, vr18, $4000
add_q   vr19, vr16, vr18
mov_l   vr20, $0
sconv_lq vr21, vr20
mul_q   vr21, vr21, $8
add_q   vr22, vr19, vr21
mov_q   vr1, vr22

```

After:

```

mov_l   vr10, $500
localaddr vr16, $0
add_q   vr19, vr16, $0
add_q   vr22, vr19, $0
mov_q   vr1, vr22

```

Before:

```

mov_l   vr13, $1
mov_l   vr10, vr13
mov_l   vr13, $11
mov_l   vr11, vr13
mov_l   vr13, $0
mov_l   vr12, vr13

```

After:

```

mov_l   vr10, $1
mov_l   vr11, $11
mov_l   vr12, $0

```

Before:

```

mov_l   vr13, $2
mov_l   vr10, vr13
mov_l   vr13, $3
mov_l   vr11, vr13
mov_l   vr13, $0
mov_l   vr12, vr13
mov_l   vr13, $1
add_l   vr14, vr10, vr13
cmpeq_l vr15, vr14, vr11
cjmp_f  vr15, .L1
mov_l   vr16, $1
mov_l   vr12, vr16

```

After:

```

cmpeq_l vr15, $3, $3
cjmp_f  vr15, .L1
mov_l   vr12, $1

```

For programs that involve a lot of variable references like example29, there aren't that many constants to pass around, so the performance tend to be very poor for that. For programs like those, especially with so many loops and redundant calculations, using LVN to remove recomputation will play a big role. But unfortunately I could not get mine working.

I am aware that the efficiency will not have a significant increase until local register allocation is implemented, because accessing register is way faster than access the cache or memory. I gave my best.

Previous versions, work history: https://github.com/YunZng/compiler_code_gen